

HPC Workshop Series: *Profiling*

Nick Featherstone
Research Computing

Slides:

https://github.com/ResearchComputing/hpc_profiling.git



Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Outline

- Basic Timing
 - the Linux “time” utility
 - timing functions in C++, Python, and Fortran
- Serial Profiling with Intel Vtune
 - Command Line
 - GUI
- Parallel Scaling Analysis

Before we begin

- ssh -X username@tlogin1.rc.colorado.edu
- ssh -X scompile
- Clone the repository (all one line):
- git clone https://github.com/ResearchComputing/hpc_profiling.git



Before we begin

- Start an interactive session:
 - `sinteractive -n4 -t60 --reservation=feathern-workshop`
 - Set up the environment and compile the code:
 - `module load intel impi python/3.5.1`
 - `cd hpc_profiling`
 - `make`



Nano survival guide

- Simple, text-based editor
- These 4 commands are useful
- Ctrl+o save (need to confirm filename)
- Ctrl+x exit
- Ctrl+k cut
- Ctrl+u paste



The Linux Time Utility

- First place to start when profiling your program

```
export OMP_NUM_THREADS=4  
time ./prog.omp (try ./prog.serial also)
```

real	0m17.801s	Wall Clock Time
user	0m58.125s	Threads x Wall Clock Time
sys	0m0.081s	System Overhead

- Works with MPI as well
- Only reports info on cpu that called mpirun

```
time mpirun -np 4 ./prog.mpi
```



Fine-Grained Timing

- Often useful to time portions of a program
- Good idea when developing your own code
- Tough when its 3rd-party software
- Useful functions:
 - Fortran: `system_clock`
 - C++ `clock()`
 - Python: `time.time()`
- Examples follow



Simple Timing (Fortran)

- Use the system_clock function (F90 onward)

```
INTEGER :: t1, t2, count_rate, count_max  
REAL*8 :: elapsed  
CALL SYSTEM_CLOCK(t1, count_rate, count_max)  
test code  
CALL SYSTEM_CLOCK(t2, count_rate, count_max)  
elapsed = real(t2-t1) / real(count_rate)  
WRITE(6,*)'Time (s): ', elapsed
```



Simple Timing (C++)

```
#include <iostream>
#include <cstdio>
#include <ctime>
int main() {
    std::clock_t tstart;
    std::clock_t t_end;
    double elapsed;
    tstart = std::clock();
    test code
    tend=std::clock();
    elapsed = ( tend- tstart ) / (double) CLOCKS_PER_SEC;
    std::cout<<"printf: "<< elapsed <<"\n";
}
```

- With c++ 11, can use std::chrono as well



Simple Timing (Python)

- Use the **time** function from the **time** module

```
import time
t1 = time.time( )
Test code
t2 = time.time( )
seconds = t2-t1
print('Elapsed time: ', seconds)
```



Advanced Timing (Python)

- Use the cProfile module to profile your code
- <https://docs.python.org/3/library/profile.html>
- Examine:
 - my_code.py
 - time_my_code.py
- Run **python time_my_code.py**



Timing with MPI

- Use the MPI_WTIME function
- Calling syntax same in C++ & FORTRAN
- Can be synchronized across all processes

USE MPI

```
REAL*8 :: t1,t2
```

```
REAL*8 :: elapsed
```

```
t1 = MPI_WTIME()
```

```
test code
```

```
t2 = MPI_WTIME()
```

```
elapsed = t2 - t1
```

```
WRITE(6,*)'Time (s): ', elapsed
```



Good Idea: Write a Timer Class

- Examine Serial_Timing.f90
- Makes it easy to change our clock function if desired
- Helps prevent typos!

Attributes

Accrued time
Time1
Time2

Methods

Start clock
Stop clock
Increment Accrued Time



Profiling with Intel Vtune

- Provided with the intel compiler
- Installed on Summit
- In-depth profiling tool
- We only have time for a quick test run
- Online tutorials here:
 - <https://software.intel.com/en-us/articles/intel-vtune-amplifier-tutorials>



Preparing to Run Vtune

- Additional environment variables need to be set:
 - Run **source ready_vtune**



Vtune: Command Line

- Runs in two stages: collection & reporting
- Have a look at serial.F90
- Try this:

```
amplxe-cl -collect hotspots ./prog.serial
```

collection

```
amplxe-cl -report top-down -r r000hs > hotspots.txt  
nano hotspots.txt
```

reporting



Vtune GUI

- Fairly sophisticated GUI interface
- Let's have a look (brace yourself)
- Follow along with me. Start by typing:
amplxe-gui



Close your interactive session

- Type exit
- Should see prompt change to shas0136 or 0137



Part II: Parallel Profiling



Outline

- What do we mean by “scaling”?
- Strong vs. Weak Scaling
- Example and Exercise

Analysis of Parallel Performance

Run a sample problem at multiple core counts and measure elapsed time.
(i.e., conduct a scaling study).

As core count increases, how does your application perform relative to ideal behavior?

Weak Scaling:

Increase problem size alongside number of cores used.

“Can I run a larger problem?”

Strong Scaling:

Fix problem size, and increase core count.

“Can I decrease my time to solution?”

Ideal Scaling

The expected performance in the absence of any communication overhead, etc.

Never achieved in practice, but standard reference point.

Ideal weak scaling:

Time-to-solution is independent of core-count (though may depend on global problem size)

Ideal strong scaling:

Time-to-solution decreases as $1/n_{\text{cores}}$

Efficiency

- Efficiency = Expected Time / Actual Time
- Expected (ideal) time is often taken relative to application's serial performance, or performance at a low core count.
- Good efficiency? Who knows? Trade-off between CPU hours and human effort.
- Suggestion: Always run your application with core counts that realize at least 80% efficiency (90 is better...).

Scaling Study Procedure


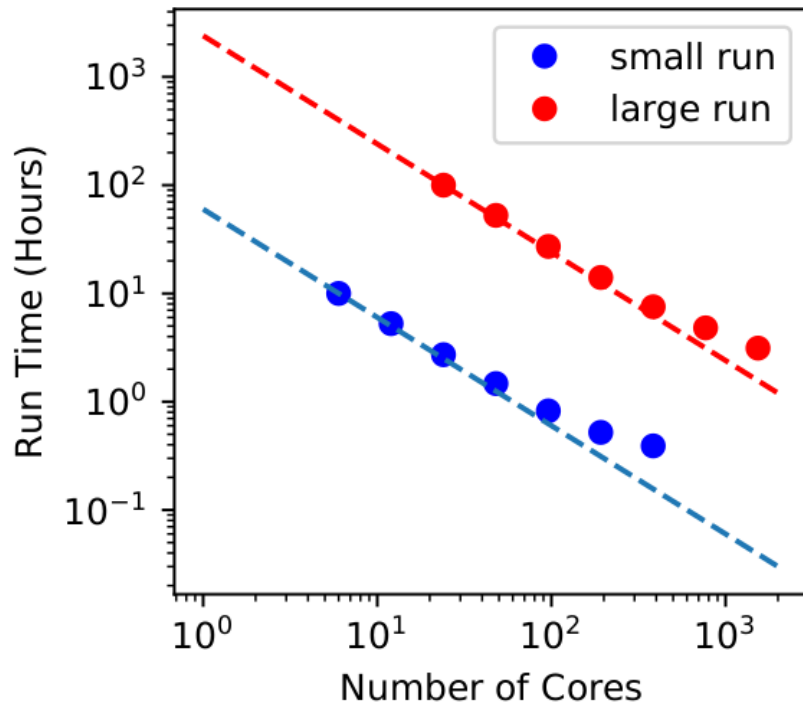
- Pick a few representative problem sizes (e.g., number of images to process, size of grid domain, number of genomes to examine)
- Decide on a sensible duration for the test. If code iterates in some fashion, run for just a few iterations.
- Run code for a short time at each problem size and with multiple core counts.
- For each test, record or calculate elapsed time, expected time, efficiency
- Plot results and ask for CPU time!

Sample Strong Scaling Data

Fixed problem size at
different core counts.

Strong-Scaling Data for Fluid Simulations			
Small-Run (128^3) Timings			
Cores	Measured Time (seconds)	Ideal Time (seconds)	Efficiency
6	60.5	60.50	1.00
12	31.84	30.25	0.95
24	16.44	15.13	0.92
48	8.90	7.56	0.85
96	4.98	3.78	0.76
192	3.15	1.89	0.60
384	2.36	0.95	0.40
Large-Run (512^3) Timings			
24	181.20	181.20	1.00
48	95.37	90.60	0.95
96	49.24	45.30	0.92
192	25.45	22.65	0.89
384	13.64	11.33	0.83
768	8.71	5.67	0.65
1536	5.6625	2.83	0.50

Strong Scaling Plots





Suggestion: Run with at least 80% efficiency

Sample Weak Scaling Data

Vary problem size in proportion to core count.

Weak-Scaling Data for Genomics Study			
64 Genomes per Core			
Cores	Measured Time (seconds)	Ideal Time (seconds)	Efficiency
6.0	60.5	60.5	1.0
12.0	61.7346938776		0.98
24.0	62.3711340206		0.97
48.0	64.3617021277		0.94
96.0	66.4835164835		0.91
192.0	68.75		0.88
384.0	70.3488372093		0.86
128 Genomes per Core			
6.0	123.0	123.0	1.0
12.0	124.242424242		0.99
24.0	125.510204082		0.98
48.0	130.85106383		0.94
96.0	139.772727273		0.88
192.0	151.851851852		0.81
384.0	157.692307692		0.78

Weak Scaling Plots



In-Class Exercise: Analyzing a Sample Program

- Examine [hpc_profiling/scale.f90](#)
- 2-D, iterative smoothing operation with nearest neighbor communication
- Grid dimensions (nx,ny) and number of iterations (nt) can be controlled via command-line arguments.
 - `mpiexec -np 16 ./prog.mpi -nx 128 -ny 256 -nt 100`
- We are going to run (but not edit) this program. We will:
 - Generate strong and weak scaling data and plot the results.

Notes

- Reusable tools provided:
 - Scaling Bash script
 - Python plotting program
- Submitting the job:
 - sbatch job.sh
- Code outputs: n_{cores} , time, nxglobal, nyglobal, niter
- A sample Python analysis script is provided.
- Let's create some scaling data together...